

# CMPE-256: Million Song Dataset Challenge\*

Carlos Hernandez  
Computer Engineering Dept.  
San Jose State University  
carlos.hernandez@sjsu.edu

Hardy Leung\*  
Computer Engineering Dept.  
San Jose State University  
kwok-shing.leung@sjsu.edu

John Lu  
Computer Engineering Dept.  
San Jose State University  
john.lu@sjsu.edu

**Abstract**—We sought to investigate the Million Song Dataset (MSD) [3], an anonymized collection of user data and metadata for popular contemporary songs. In the accompanying Kaggle MSD Challenge (MSDC), over 150 teams proposed solutions to predict which songs a user will listen to, given the collective listening history of all users. The sizable training data contains over 1M listeners, 380K songs, and 48M user-song-play triplets. In this project, we studied multiple collaborative filtering techniques; our best result of 17.453% would have placed us in the second place in the Kaggle Challenge.

**Index Terms**—recommender system, implicit feedback

## I. INTRODUCTION

The Million Song Dataset (MSD) is a collection of user data and metadata for popular contemporary songs. MSD was designed to encourage research on algorithms that scale to commercial sizes, and to provide a reference dataset for evaluating research [3]. Most relevant to our project is the taste profile of over 1M listeners, captured in 48M user-song-play triplets. While users are fully anonymized, MSD provides a rich collection of information related to the songs, including genre classification from the Tagma dataset [10], the MusiXmatch lyrics dataset [3], and even acoustic and spectral analysis. In its entirety, the MSD contains a total of 280GB worth of data [3].

In the MSD Challenge, each of the 48M user-song-plays triplets is a positive integer representing the number of times a user has played a specific song. While we will adopt familiar terminologies as item-based (IB) and user-based (UB) collaborative filtering, and matrix factorization (MF), we are nonetheless making recommendations based on *implicit* feedback, where the non-zero entries in the utility matrix should be interpreted as merely a signal, rather than a definite opinion or rating.

## II. RELATED WORK

The Netflix Prize competition that began in 2006 was a key driving force behind the rapid adoption of many collaborative filtering techniques across high-tech industries [2], such as the popularization of matrix factorization

by Simon Funk in his tongue-in-cheek blog post [4]. However, it is rare to find datasets that employ explicit ratings in real-life. More frequently, the feedback would come implicitly in the form of a click, a listen, a purchase, and so on. While a positive feedback does signal preference, an absence of feedback should not be interpreted as a definitive negative opinion. In [5], Hu et al. adopted a formulation where implicit feedbacks are coupled with *confidence*, modeled as larger weights in the objective function; they also made a clever and necessary adaptation of the Alternating Least Square (ALS) algorithm to handle the dense loss function induced by their formulation, where even empty terms receive a non-zero, albeit diminutive, penalty. Subsequently, Johnson [7] extended Hu’s work with a logistic formulation that directly models the probability that a user will prefer a specific item.

In the MSD Challenge, the objective is to predict the song preference of each of the 110K users in the testset, given half of their song preference, and the complete listening history of 1M users in the trainset. We were to offer an individualized and ranked list of  $\tau$  recommended songs to each of those users, where  $\tau = 500$  in the MSD Challenge. The prediction is then measured with the truncated mAP metric. We first need to define *precision@k* to be the fraction of correctly predicted songs within the first  $k$  recommendations. Then, the truncated mAP metric evaluates *precision@k* at each rank with a correct prediction, divided by the number of the hidden songs (correctly predicted or not). For example, if a user  $u$  has 5 hidden songs, 3 of which were ranked within the 500 predictions at 3<sup>rd</sup>, 7<sup>th</sup>, and 20<sup>th</sup> places, and the remaining two were not predicted at all. Then, the truncated mAP for the user is:

$$\frac{1}{5} \left( \frac{1}{3} + \frac{2}{7} + \frac{3}{20} \right) \simeq 0.15381$$

This example shows that we can achieve a higher mAP score by correctly predicting more songs at higher ranks.

## III. OUR APPROACH

Armed with a good understanding of what the project entails, we set out to experiment with the following

\* <https://github.com/jfantab/cmpe256-project>

techniques, each of which was carefully considered beyond simple application of existing packages.

- algorithm design and architecture engineering
- neighborhood-based collaborative filtering
- matrix factorization for implicit feedback
- similarity based on factorized latents
- similarity based on embeddings of metadata
- term-frequency document-frequency normalization
- ensemble recommendation

#### A. Algorithm Design and Architecture Engineering

A crucial decision was made to develop the program using C++, and most of the experiments were run on a modestly equipped laptop. We did evaluate a few options, including HPC, Google CoLab, and various Python and Java packages. At the end, we decided that having a flexible and properly engineered architecture is just as important as choosing the right algorithms, even more so if the algorithms were designed with the architecture in mind. The C++ code was written completely from scratch, with the exception of a small JSON parser, and thus we avoided significant amount of external dependencies, such as GPU availability and reliability, timeout issues with cloud-based computing, and lack of transparency in implementation characteristics of gray-box packages. We engineered the algorithms to take full advantage of multi-threaded parallelism, and to fit under approximately 4GB of RAM to avoid disk thrashing. While this was an old-school approach to engineering and optimization, doing so allowed us to run various algorithms in a reasonable amount of time. We did make multiple design tradeoffs given several naïve implementations would have required  $O(|U| \times |I|)$ ,  $O(|U|^2)$ , or  $O(|I|^2)$  memory, where  $U$  is the set of users, and  $I$  is the set of songs. Since  $|U| \simeq 1.1M$  and  $|I| \simeq 380K$ , those implementations would have simply run out of memory.

Our basic architecture is simple yet designed with performance in mind. The utility matrix  $M$  is implemented as two complementary matrices called UIR and IUR (which stand for User-Item-Rating and Item-User-Rating), implemented as indexed arrays to afford  $O(1)$  lookup of both  $I_u$  (all items rated by the user  $u$ ) and  $U_i$  (all users who rated the item  $i$ ). UIR is sorted first by  $U$  and then by  $I$ , whereas IUR is sorted first by  $I$  and then  $U$ . This allows us to enumerate the set  $U_i$  in increasing order of user, and similarly for  $I_u$ . The memory requirements of UIR and IUR are both  $|M|$ . In our case,  $|M| \simeq 48$  million. Note that  $\sum_{u \in U} |I_u| = \sum_{i \in I} |U_i| = |M|$ .

The simplest approaches to solving MSD would be the familiar user-based (UB) and item-based (IB) collaborative filtering, which we implemented first as a baseline,

and then carefully extended. Note that the number of plays should not be used as rating, since it is both unreliable and unbounded. Moreover, since our goal is to predict *whether* a user would play a song, it makes sense to binarize the utility matrix  $M$  so that an entry is 1 if the user played the song at least once, and 0 otherwise. In the following, we assume  $M$  has been binarized unless stated otherwise.

While UB and IB are conceptually simple, we have two serious problems. First,  $U$  and  $I$  are both very large. Second, we need to make a total of  $|U_{\text{test}}| \times |I|$  predictions, where  $U_{\text{test}}$  is the subset of users to receive our top-500 recommendations<sup>1</sup>. In IB, we'll look at each item  $i \in I_u$ , and compute  $\text{sim}(i, j)$ , defined to be the similarity score between items  $i$  and  $j$ . Even worse,  $\text{sim}(i, j)$ , if again evaluated individually, would require going through both  $U_i$  and  $U_j$ , taking at least  $O(|U_i| + |U_j|)$  time. All these add up to a serious complexity problem.

This problem is not unique to us. In fact, the Python Surprise package [6] precomputes an  $O(|I|^2)$  array of  $\text{sim}(i, j)$ , since we need to access the value of  $\text{sim}(i, j)$  very frequently. The problem is that  $O(|I|^2)$  is a very large number ( $\simeq 144$  trillion entries in our case), and Surprise would quickly run out of memory. Clearly, we need a different approach.

We still want to compute all  $|I|^2$  pairs of  $\text{sim}(i, j)$ , and we process each item  $i$  one at a time, but all  $\text{sim}(i, \star)$  would be computed simultaneously. To address the memory problem, for each item  $i$ , we would keep at most  $K$  best (i.e. largest)  $\text{sim}(i, j)$ .

For each  $i$ , we first initialize a  $\text{count}(i, \star)$  array to zero, and go through each user  $u \in U_i$ , and for each item  $j$  in  $I_u$ , we simply increment  $\text{count}(i, j)$ . This simple sequential pass allows us to count, for each  $j$ , the number of users who played both  $i$  and  $j$  (keep in mind that we are working on a particular item  $i$ ). Next, we turn  $\text{count}(i, \star)$  into  $\text{sim}(i, \star)$  with a straight-forward formula in  $O(1)$  time. Finally, we will find the collection of  $K$  largest  $\text{sim}(i, \star)$  using a linear-time median-finding algorithm<sup>2</sup>, and we store the result into an array of size  $K$ . In our implementation, we set  $K = 500$ .

Together, computing all  $\text{sim}(i, \star)$  takes  $(|I| + |M_{U_i}|)$  time, where  $M_{U_i}$  is the subset of the utility matrix  $M$  taking only rows corresponding to users in  $U_i$ . Summing

<sup>1</sup>To put things in perspective,  $|U_{\text{test}}| \times |I| = 42$  trillion, but the Netflix Prize only asked for 2.8 million predictions.

<sup>2</sup>We used `std::nth_element()`, which found not only the  $n^{\text{th}}$  element, but all elements smaller than said element. This is done in linear time in the size of the vector.

over all possible  $i$ , the total runtime complexity of our IB implementation is:

$$\sum_{i \in I} O(|I| + |M_{U_i}|) = O(|I|^2) + O(\sum_{i \in I} |M_{U_i}|)$$

which can be rewritten as:

$$O(|I|^2) + O(\sum_{u \in U} |I_u|^2) = O(|I|^2 + \beta_{\text{item}}^M |M|)$$

where  $\beta_{\text{item}}^M$  is a (non-constant) multiplier determined by the characteristics of the matrix  $M$ . It is as if we would process all items of user  $u$ , namely  $I_u$ , exactly  $|I_u|$  times. This does imply a quadratic behavior, albeit on a per-user basis. Empirically,  $\beta_{\text{item}}^M = 118$ . Note that  $\beta_{\text{item}}^M = O(|I|)$  but in practice  $\beta_{\text{item}}^M$  is likely significantly smaller than  $|I|$  (by 3200X).

Let  $\text{sim}^K(i, j)$  denote the abridged version of  $\text{sim}(i, j)$ , where only the best  $K$  entries of  $\text{sim}(i, \star)$  are kept for each  $i$ . The memory requirement is therefore  $O(K|I|)$  for IB and similarly  $O(K|U|)$  for UB. In our case, given  $K = 500$ , they amount to 190M and 555M entries respectively for IB and UB collaborative filtering. We can further improve the memory footprint by reducing  $K$ , but the entirety of  $\text{sim}^K(i, j)$  needs to be kept in memory for downstream usage. We did perform a cursory analysis on the impact of reducing  $K$  on the performance of our IB recommender, as shown in Figure 1 (note that the  $x$ -axis is on a logarithmic scale, and  $500 \simeq 2^9$ ). It seems reducing  $K$  to  $64 = 2^6$  or even  $32 = 2^5$  has negligible impact on performance.

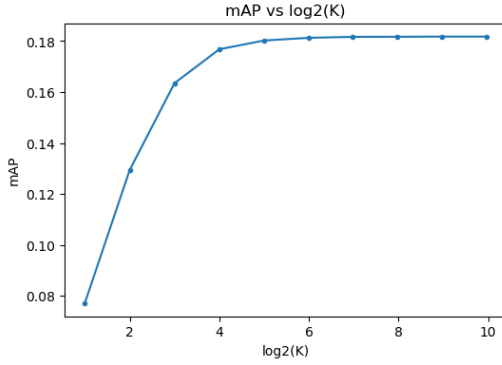


Fig. 1: The mAP performance of an IB run with different values of  $K$ , where  $K$  is the maximum number of  $\text{sim}(i, \star)$  kept in memory for each item  $i$ .

Once we conclude the computation of  $\text{sim}^K(i, j)$ , we perform recommendation for each user, starting by initializing a score array of size  $|I|$  to zero.  $\text{rec}(u, i)$  represents the prediction (or recommendation) of item  $i$  to user  $u$ . We will go through each item  $i$  in  $I_u$ , retrieving all  $K$  items in  $\text{sim}^K(i, \star)$ , and increment  $\text{rec}(u, j)$  as follows:

$$\text{rec}(u, j) \leftarrow \text{rec}(u, j) + (r_{u,i}) (\text{sim}(i, j))$$

for each eligible  $j \in \text{sim}^K(i, \star)$ . In essence, we concurrently update all  $\text{rec}(u, \star)$ . The update per user takes  $O(K \times |I_u|)$  time, and thus in aggregate the complexity for all  $U_{\text{test}}$  users is  $O(K \times |M_{\text{test}}|)$ . Since  $K = 500$  and  $|M_{\text{test}}| \simeq 0.1 \times |M|$ , it would be similar to going through the utility matrix  $M$  50 times.

Note that the computation of both  $\text{sim}(i, j)$  and  $\text{rec}(u, i)$  can be parallelized. We decided that a multi-threaded approach would perform best, and in fact were able to achieve near linear speedup, close to 750% CPU utilization with a hyperthreaded Quad-Core CPU. To reiterative, the runtime complexity of our IB implementation is:

$$O(|I|^2 + \beta_{\text{item}}^M |M|)$$

For our particular example,  $|I|^2$  and  $\beta_{\text{item}}^M |M|$  evaluate to  $144 \times 10^9$  and  $6 \times 10^9$  respectively (though we caution that we should not add these two numbers). While it still takes significant amount of time to run the IB and UB algorithms, whose runtimes were dominated by the  $\text{sim}(i, j)$  and  $\text{sim}(u, v)$  computation, these numbers are not too daunting for a well-polished algorithm running on a modestly equipped laptop.

We can derive a similar analysis for UB calculation:

$$O(|U|^2) + O(\sum_{i \in I} |U_i|^2) = O(|U|^2 + \beta_{\text{user}}^M |M|)$$

where  $\beta_{\text{user}}^M = 4445$ . Overall, the IB and UB runtimes were dominated by  $O(|I|^2)$  and  $O(|U|^2)$  respectively.

## B. Neighborhood-based Collaborative Filtering

We spent a significant amount of time on infrastructure engineering, and the investment started to pay off when we can afford to reason about IB and UB approaches more deeply. We start by acknowledging the work of Aioli [1] who presented an elegant generalization of cosine similarity for neighborhood-based collaborative filtering. Instead of the standard formulation for binary values:

$$\text{sim}_{\text{std}}(i, j) = \frac{|U_{i,j}|}{|U_i|^{1/2} \times |U_j|^{1/2}}$$

where  $U_{i,j}$  is the set of users who rated both items  $i$  and  $j$ . He proposed a variant that captures the asymmetry between  $i$  and  $j$ , in that we use  $\text{sim}(i, j)$  to predict whether the user will (or will not) like item  $i$ , while it is already *known* that the user likes item  $j$ . As such, he proposed to generalize the above equation<sup>3</sup>: to:

$$\text{sim}_\alpha(i, j) = \frac{|U_{i,j}|}{|U_i|^{1-\alpha} \times |U_j|^\alpha}$$

<sup>3</sup>Note that we adopted the opposite definition of  $\alpha$  compared to Aioli. Due to time constraint, we chose to only disclose such discrepancy rather than to resolve the inconsistency.

where  $\alpha$  is a tunable hyper-parameter. In fact, from a Bayesian perspective,  $\text{sim}_\alpha(i, j)$  is equivalent to  $P(j|i)^{1-\alpha} \times P(i|j)^\alpha$ . When  $\alpha = 0.5$ ,  $\text{sim}_\alpha(i, j)$  reduces to cosine similarity. When  $\alpha = 1$ , this reduces to a simple conditional probability of seeing  $i$  when  $j$  has been observed. Varying the value of  $\alpha$  allows us to trade off between simple alignment (whether  $i$  and  $j$  are similar) and prediction (whether seeing  $j$  tells us anything about  $i$ ).

Another important concept introduced by Aioli is *locality* [1]. The idea is that highly similar items should be disproportionally weighted. In other words, instead of calculating the score of an item  $i$  with a weighted sum:

$$\sum_{j \in I_u \setminus \{i\}} (\text{sim}_\alpha(i, j)) \times r_{u,i}$$

we could use an exponentiated version of the weights:

$$\sum_{j \in I_u \setminus \{i\}} (\text{sim}_\alpha(i, j))^\gamma \times r_{u,i}$$

If  $\gamma > 1$  we put more emphasis on higher similarity scores, and vice versa for  $\gamma < 1$ . When  $\gamma = 0$ , each term reduces to 1, and hence the recommendation simplifies strictly to a popularity contest.

### C. Matrix Factorization for Implicit Feedback

Since the user-song-play triplets in the MSD dataset is inherently an implicit feedback, we sought to apply techniques that were designed for this purpose, including the work by Hu et al. [5] who introduced the concept of preference and confidence. While the former is just the prediction whether the user would play the song, the *confidence* is how much emphasis we put on matching this preference. In their RMSE-like loss function, the penalty is  $c_{u,i}(p_{u,i} - \hat{p}_{u,i})^2$ , where  $c_{u,i}$  is either 40 when the preference is positive, or 1 when it is negative. In a related work, Johnson [7] proposed a maximum-likelihood formulation solvable using logistic regression, but with the same concepts of confidence and preference. Instead of reinventing the wheels, we adopt Frederickson’s implementation in the Python package Implicit<sup>4</sup>. Both works adapted the matrix factorization formulation to handle implicit feedback, and similarly generated both a  $|U| \times d$  user matrix  $U_{\text{emb}}$ , and a  $|I| \times d$  item matrix, where  $d$  is the chosen number of embedding dimensions (we take the defaults of 128 dimensions for Hu’s method (referred to as ALS), and 32 dimensions for Johnson’s method (referred to as LMF). These embeddings can then be used to compute any  $r_{u,i}$  by taking the inner product of the corresponding row and column. We can then collect and sort  $r_{u,\star}$  for each user in  $U_{\text{test}}$ , and recommend the top  $\tau$  songs ( $\tau = 500$

for our purpose). Given the matrices  $U_{\text{emb}}$  and  $I_{\text{emb}}$ , the runtime complexity of our recommendations is a modest  $O(|U_{\text{test}}| \times |I| \times d)$ .

### D. Similarity based on Factorized Latents

We experimented with using embeddings obtained from matrix factorization directly in neighborhood-based recommendation. In other words, instead of calculating  $\text{sim}_\alpha(i, j)$  based on ratings, what if we compute  $\text{sim}_{\text{emb}}(i, j)$  by taking inner product of the two  $d$ -dimensional embeddings  $I_{\text{emb}}(i)$  and  $I_{\text{emb}}(j)$ ? We attempted to do that, using a similar methodology described earlier to pre-compute the top  $K$   $\text{sim}_{\text{emb}}(i, \star)$  for each item  $i$ . Since we do compute all possible  $\text{sim}_{\text{emb}}(i, \star)$ , the runtime complexity is  $O(d \times |I|^2)$  since it takes  $O(d)$  time to perform a single dot product. This is significantly slower than our item-based approach described earlier. Notably, since  $d$  is either 128 (ALS) or 32 (LMF), this approach is more than one or two orders of magnitude slower. We were however curious whether this method combines the best of neighborhood-based and matrix factorization.

### E. Similarity based on Embeddings of Metadata

One of the very first ideas we brainstormed was to investigate how to use the vast collection of metadata made available to us. Due to time constraint, we initially focused on possibly the most interesting part of the metadata – lyrics [3]. Our first task was to properly extract the information of each song from the Muxmatch lyrics files. The Muxmatch dataset only contains data for 237,680 songs, or 62% of the total number of songs in the Taste Profile subset. According to the authors, there were a variety of reasons why data is incomplete, such as copyright restrictions. Moreover, the words were organized in a bag-of-words format in order to avoid any copyright infringement claim, which limited our options in performing more meaningful semantics analysis.

For each song with lyrics, we proceeded to parse and extract the terms and their counts into a SCIPY CSR sparse matrix format, and further applied TF-IDF [11] transformation to normalize the term frequencies (so that frequently mentioned words due to repeated verses do not overwhelm the analysis). In the process, we handled a variety of issues, including conversion between stemmed and unstemmed words, non-English words, and various debugging issues related to GLOVE [8]. Once properly cleaned, and TF-IDF transformed, we obtained a single embedding of the lyrics as a sum of the embeddings of the words, weighted by the term frequencies.

These embeddings were then used as latent vectors for each item in exactly the same fashion as we did in

<sup>4</sup><https://github.com/benfred/implicit>

the previous subsection. Similarly, the complexity is  $O(d \times |I|^2)$  for item-based collaborative filtering using the lyrics embeddings, where  $d = 50$ .

#### F. Term-Frequency Document-Frequency Normalization

Earlier we mentioned our analysis is based on a binarized utility matrix, as the play count data was deemed too unreliable [1]. Let's say if a certain listener has played a song 2213 times<sup>5</sup>, how much does she like the song compared to the ones she played 167 times, 30 times, 3 times, or only once? Indeed, once again we adopted an idea from information retrieval that was also suggested by Frederickson in his Implicit package, namely that the user-song utility matrix can be treated as a term-document matrix, where a song is a document, and a user is a word (or term). Using the above example, this special user is analogous to a term that appears too many times in certain documents. Therefore, we took a clue from information retrieval and tried applying TF-IDF [11] and the BM25 improvement [9]. The advantage of these normalization techniques is that they do not reduce the sparsity of the utility matrix, as empty entries stay empty.

Our earlier analysis can be easily adapted to handle numerical entries transformed by TF-IDF or BM25 other than just binary values – the same (raw) cosine similarity works in a similar fashion. Unfortunately, due to an implementation bug, we were not able to confirm whether TF-IDF or BM-25 cosine similarity yielded better or worse results compared to the binary version.

On a related note, we have also considered the use of Jaccard similarity, which is similar to cosine similarity with binary values. In particular,

$$\text{sim}_{\text{Jaccard}}(i, j) = \frac{|U_{i,j}|}{|U_i| + |U_j| - |U_{i,j}|}$$

We did not perform a full-blown analysis, but anecdotal evidence suggested that  $\text{sim}_\alpha(i, j)$  performs better and is more versatile given the value  $\alpha$  can be tuned. In all of our subsequent discussion, we exclusively use  $\text{sim}_\alpha(i, j)$  on the binarized utility matrix.

#### G. Ensemble Recommendation

We have access to a wide variety of recommenders, from IB and UB collaborative filtering, matrix factorization, embedding-based neighborhood analysis, even the naïve popularity-based method. Each recommender is able to make some positive recommendations. Can we combine them to achieve better results? The answer is most definitely yes. In our project, we evaluated two different ensemblers  $E_A$  and  $E_B$ .

<sup>5</sup>This person actually exists in the dataset.

We shall start with a set of recommenders  $R$ , each recommending a ranked list of  $\tau$  songs for each user. In our first ensembler  $E_A$ , we would assign a weight  $w_r$  to each recommender  $r$  using a rough estimate of its standalone mAP based on past training. We then estimate the probability  $p_A(r, k)$  that the recommender  $r$  made a correct guess at rank  $k$ . A very crude estimate would be  $\frac{1}{k}$  (i.e. 100% certainty at the first recommendation, 25% certainty at the fourth, and 5% certainty at the twentieth). We'll then multiply the probability  $p_A(r, k)$  with  $(w_r)^\rho$  to form a score for each ranked song, where  $\rho$  is a tunable hyper-parameter. As a result, we now have  $|R| \times \tau$  songs recommended by the  $|R|$  recommenders, with some songs appearing multiple times. The score represents the relative confidence of the recommendation adjusted by the strength of the recommender. We will then sort and pick the best  $\tau$  unique songs from the sorted pool of recommended songs.

In our second ensembler  $E_B$ , we will once again collect the  $|R| \times \tau$  songs. For each unique song in the list, we will find the ranking of the song from each of the recommenders, and based on the ranking of the song, and the strength of the recommender, we compute the probability that it is a correct guess. For example, we could perform a recommender-specific polynomial fit on the historical *precision@k* data (see Figure 2) to estimate the probability of a correct prediction. If a recommender did not rank a song, its rank is assumed to be  $\tau$ . As such,  $p_B(r, i)$  takes on the meaning of a true probability. We assume further that these probabilities are independent, and use them to compute the probability that the song  $i$  is indeed a correct guess as follows:

$$p_B(i) = 1 - \left( \prod_{r \in R} (1 - p_B(r, i)) \right)$$

The idea is that, supposed we have good estimate for the probabilities of a correct guess on the given song, and those probabilities are independent (an obviously wrong assumption, which we made anyway), then the probability that the song is indeed a correct prediction can be computed with simple probability formulas. We will then resort and select the top  $\tau$  songs according to these probabilities.

## IV. EXPERIMENTAL RESULTS

In this section, we shall present the details of our experiments. The Kaggle MSD Challenge dataset contains a utility matrix  $M_{\text{train}}$  describing the listening habit of all fully visible users in the trainset. We are also given




Fig. 2: Average  $precision@k$  for IB, UB, and popularity-based recommenders.

$M_{test}^V$  and  $M_{test}^H$ , the visible and hidden halves of the song history of the users in  $U_{test}$ <sup>6</sup>.

We merged  $M_{train}$  and  $M_{test}^V$  into  $M_{all}$ . Next, we randomly carved out 10% of the non-test users to form  $U_{valid}$ , and randomly hide half of these entries in the utility matrix for the purpose of validation. This methodology allows us to test and tune our recommenders without peeking at the ground truth. We then perform various experiments with different algorithms. The baseline result is the popularity-based recommender, with an mAP score of only 2.3%.

#### A. Neighborhood-based Recommenders

Thanks in no small part to our architectural investment, we were able to perform analysis on various combinations of hyper-parameters despite the tight schedule. On average, individual IB and UB runs took approximately 2 to 3 minutes each,

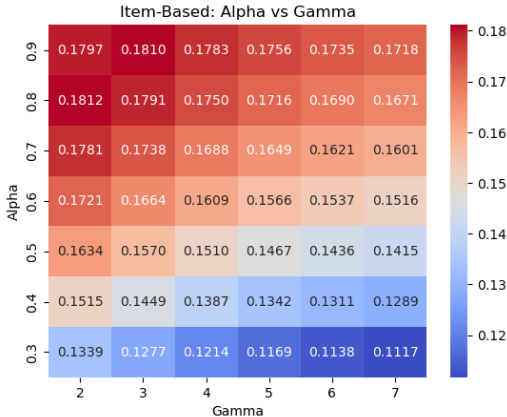


Fig. 3: Grid Search on  $\alpha$  and  $\gamma$  for IB recommender.

<sup>6</sup>During the competition, the contestants did not have access to  $M_{test}^H$ , but  $M_{test}^H$  was released afterwards for posterity. We evaluated our recommender only after all hyper-parameter tuning was done.




Fig. 4: Grid Search on  $\alpha$  and  $\gamma$  for UB recommender.

$\alpha$  and  $\gamma$  are by far the most important hyper-parameters that determine the performance of our IB and UB recommenders. Figures 3 and 4 show how mAP varies as a function of  $\alpha$  and  $\gamma$ . We further performed a local grid search around the hottest spots, and founded the best parameters to be  $\alpha_{item} = 0.85$  and  $\gamma_{item} = 2.5$  for item-based recommender, and  $\alpha_{user} = 0.65$  and  $\gamma_{user} = 4.0$  for user-based recommender. The mAP for the IB recommender alone is 18.18%, and that for the UB recommender alone is 15.38%.

#### B. Ensemble Recommender

Our ensemble recommender has a tunable parameter  $\rho$ , which modifies the weights of the input recommenders. In the experiment shown in Figure 5, we made a recommendation based on an IB recommender and a UB recommender, and varied  $\rho$  between 0.25 and 5.00 inclusive. The weights for the two recommenders were simply 0.18 and 0.15 respectively. The ensemble recommender outperforms both IB and UB recommenders across all choices of  $\rho$ . At the best value of  $\rho$ , the mAP of the ensembler exceeds that of IB by 2.3%, and that of UB by 20.9%.

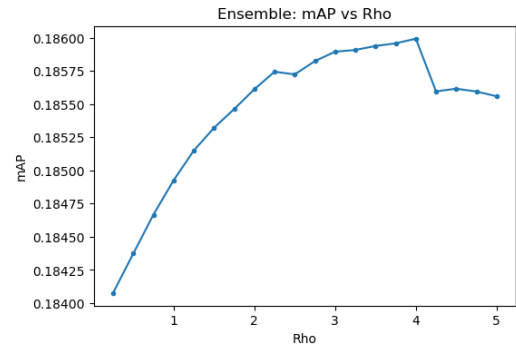


Fig. 5: Grid Search on  $\rho$  for ensemble recommender.

### C. Comparison on the Validation Set

We now compare the mAP and runtime performance of different algorithms. Performance is measured with both the truncated mAP metric evaluated at  $\tau = 500$  (see [3]), and a metric called coverage, which simply measures the average percentage of songs in the ground truth correctly predicted by the algorithm. All experiments were performed on a single Quad-Core 2.33GHz i5 MacBook Pro circa 2018 with 8GB RAM, and all numbers quoted were measured with a separate validation testset described earlier. ITEM-BASED was run with parameters  $\alpha_{\text{item}} = 0.85$  and  $q_{\text{item}} = 2.5$ , and USER-BASED was run with parameters  $\alpha_{\text{item}} = 0.65$  and  $q_{\text{item}} = 4.0$ . The complete ENSEMBLER algorithm utilized the output of these two recommenders, with  $\rho = 4.0$ . We also included both the raw runtime of ENSEMBLER and those of ITEM-BASED and USER-BASED. All in all, ENSEMBLER performed the best all other algorithms, achieving an mAP score of 18.599%.

The first and most obvious observation is the poor performance of *all* embedding-based algorithms across the board, whether it is from lyrics embeddings, matrix factorization including ALS [5] and LMF [7], or using those embeddings in a neighborhood-based setting. In fact, because of such poor performance, we spent significant amount of time looking for bugs.

The second observation is that the “KISS” principle is readily applicable. One would expect the more information we utilize, or the more sophisticated the algorithms, the better result we would get. The reality is more complicated. For example, we tried multiple ways to transform the user-song-play triplets into more usable form, including TF-IDF and BM-25. At the end, a better understanding and more elegant generalization of binary cosine similarity led to better performance. In the case of the ensemble algorithm, the more sophisticated and theoretically sound ensembler  $E_B$  lost out to the simpler ensembler  $E_A$ .

### D. A Moment of Truth with Kaggle

Unfortunately, despite the great performance of our ensemble algorithm, it did not perform as well as we expected when we run it on the actual testset. We achieved an mAP score of 17.453%, which is a significant drop from our best result of 18.599% evaluated on the validation set. In retrospect, we realized we made a major mistake in not conducting enough cross-validation during grid search, and hence the result might have been somewhat overfitted. Sadly, we did not have enough time to correct course. That said, the mAP score of 17.453% would have placed our team in second place in the Kaggle leaderboard, behind only the winner at 17.909%.

### V. DISCUSSION

While disappointing at first, we realized the goal of the project has never been to achieve the top ranking. Instead, we gained an appreciation for good architecture and implementation. On the other hand, while it is possible that what we saw was the result of an extremely embarrassing bug, we may have an explanation for the surprisingly poor performance of embedding-based approaches.

Let’s first take a detour to revisit the Netflix Prize, in which the objective was to minimize the RMSE of rating prediction. It would be three long years before the winning team beat the prior art by 10%, achieving an RMSE of 0.8563. Since the lowest and highest movie ratings are 1 and 5 respectively, the RMSE spans a whopping  $\frac{0.8563}{5-1} \simeq 21.4\%$  of the ranges. These RMSE represents a combination of inherent noise and uncertainty that even the best recommender cannot overcome.

In the MSD Challenge, on the other hand, the ability to correctly predict the items early in the ranking would make or break the recommenders. Here lies an important difference between RMSE (Netflix) and mAP (MSDC): while RMSE emphasizes on accuracy in an equitable fashion (accurate ratings on bad movies are valued just as much as those on good movies), mAP focuses exclusively on the very top rankings. In our dataset, the average number of items per user is  $\frac{|M|}{|U|} \simeq 48$ , a miniscule number comparing to the  $380K$  items. From this perspective, even the 2.3% mAP achieved by the popularity-based recommender looks impressive!

But how come embedding-based recommenders perform so much worse than IB and UB recommenders? The key may be the *quality* of the similarity measures. In Figure 6, we took a snapshot of the  $\text{sim}(i, \star)$  vector each time we process a new item  $i$ , sorted them and normalized them so the largest value is 1. Afterwards, we sampled and tallied each half-percentage between 0% and 10%. In other words, we are interested in the dropoff in the values of similarity across all items. Notice how the similarity scores decrease at different rates.

We postulate that the structure of the neighborhood landscape, as measured by similarity, plays a significant role in the accuracy of prediction. It may not be very helpful to have too many similar items, just like it would be annoying and unproductive to hear “most anything is good here!” when asking the waitress to recommend a dish. We have consistently observed that steeper curves coincide with better mAP performances. In fact, this would explain why increasing the locality factor  $\gamma$  improves performance, as it is an inexpensive way to steepen the curve. Perhaps this is why in the



ALGORITHM	PARAMETERS	MAP	COVERAGE	WALL RUNTIME (SEC)	CPU RUNTIME SEC
RANDOMIZED		0.001%	0.122%	0	0
EMBEDDING		0.016%	0.666%	4278	31331
LMF [7]		1.586%	20.399%	326	2302
POPULARITY		2.354%	18.887%	0	0
ALS [5]		4.671%	24.081%	1391	9147
USER-BASED (UB)	$\alpha(0.65), q(4.0)$	15.376%	49.007%	159	1147
ITEM-BASED (IB)	$\alpha(0.85), q(2.5)$	18.176%	57.546%	116	815
ENSEMBLER on (IB, UB)	$\rho(4.0)$	18.599%	58.385%	13 (288)	13 (1975)

TABLE I: Comparison among various algorithms. All runtimes exclude the overhead of reading the input data and minor preprocessing. The mAP metric reports the truncated version at  $\tau = 500$  (see [3]). The COVERAGE metric measures the average percentage of songs in the ground truth correctly predicted by the algorithm. All numbers quoted were measured with a validation test set randomly splitted from the training data. The ENSEMBLER runtime in parenthesis includes those of the item-based and user-based runtime.

grid-search we found  $\gamma_{\text{user}} > \gamma_{\text{item}} > 1$ .

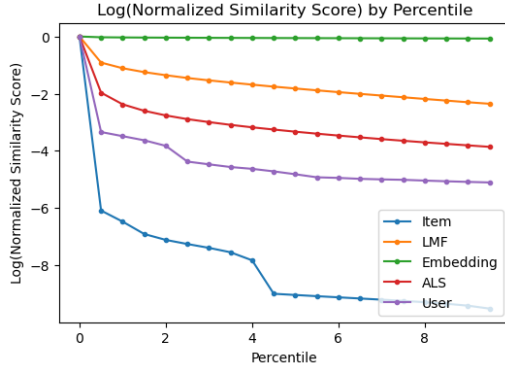


Fig. 6: Logarithm of normalized similarity score at each precision point ( $\text{precision}@k$ ) for each recommender.

In conclusion, we have investigated the Million Song Dataset Challenge and implemented multiple algorithms. Our solution was built on a robust, efficient, and flexible architecture that allowed us to explore and fine-tune various possibilities.

#### APPENDIX – WORK ASSIGNMENT

The initial intent was that all three members make equal (or at least equitable) contributions towards the project.

We share research and background discovery equally among Carlos, Hardy, and John. A significant amount of time was spent on nailing down the project choice and refinement.

Hardy was responsible for all C++ implementation including collaborative filtering, ensembler, and the use of Python Implicit (ALS and LMF). John was responsible for the development of embedding assignment including a significant amount of research on how to create useful embeddings from WORD2VEC, and GLOVE. John also

looked into other topics related to embedding such as genre classification.

The report was written by Hardy, including various experiments that accompanied the analysis (e.g. Discussion section). John made the presentation.

#### REFERENCES

- [1] AIOLLI, F. Efficient top-n recommendation for very large scale binary rated datasets. In *Proceedings of the 7th ACM conference on Recommender systems* (2013), pp. 273–280.
- [2] BENNETT, J., LANNING, S., ET AL. The netflix prize. In *Proceedings of KDD cup and workshop* (2007), vol. 2007, New York, p. 35.
- [3] BERTIN-MAHIEUX, T., ELLIS, D. P., WHITMAN, B., AND LAMERE, P. The million song dataset. In *Proceedings of the 12th International Conference on Music Information Retrieval (ISMIR 2011)* (2011).
- [4] FUNK, S. Try this at home. <http://sifter.org/~simon/journal/2006> (2006).
- [5] HU, Y., KOREN, Y., AND VOLINSKY, C. Collaborative filtering for implicit feedback datasets. In *2008 Eighth IEEE international conference on data mining* (2008), Ieee, pp. 263–272.
- [6] HUG, N. Surprise: A python library for recommender systems. *Journal of Open Source Software* 5, 52 (2020), 2174.
- [7] JOHNSON, C. C. Logistic matrix factorization for implicit feedback data. *Advances in Neural Information Processing Systems* 27, 78 (2014), 1–9.
- [8] PENNINGTON, J., SOCHER, R., AND MANNING, C. D. Glove: Global vectors for word representation. In *Proceedings of the 2014 conference on empirical methods in natural language processing (EMNLP)* (2014), pp. 1532–1543.
- [9] ROBERTSON, S. E., WALKER, S., JONES, S., HANCOCK-BEAULIEU, M. M., GATFORD, M., ET AL. Okapi at trec-3. *Nist Special Publication Sp 109* (1995), 109.
- [10] SCHREIBER, H. Improving genre annotations for the million song dataset. In *ISMIR* (2015), pp. 241–247.
- [11] SPARCK JONES, K. A statistical interpretation of term specificity and its application in retrieval. *Journal of documentation* 28, 1 (1972), 11–21.